

## Data loading

```
In [4]: 1 import pandas as pd
        2 age_train=pd.read_csv("age_train.csv",names=["uId", "age_group"])
        3 age_train.tail(10)
```

Out[4]:

|         | uld     | age_group |
|---------|---------|-----------|
| 2009990 | 3512491 | 3         |
| 2009991 | 3512492 | 2         |
| 2009992 | 3512493 | 3         |
| 2009993 | 3512494 | 2         |
| 2009994 | 3512495 | 5         |
| 2009995 | 3512496 | 1         |
| 2009996 | 3512497 | 6         |
| 2009997 | 3512498 | 2         |
| 2009998 | 3512499 | 2         |
| 2009999 | 3512500 | 3         |

```
In [5]: 1 age_test=pd.read_csv("age_test.csv",names=["uId"])
        2 age_test.tail()
```

Out[5]:

|        | uld     |
|--------|---------|
| 502495 | 3180767 |
| 502496 | 3181007 |
| 502497 | 3181086 |
| 502498 | 3181099 |
| 502499 | 3181144 |

```
In [6]: 1 user_basic_info=pd.read_csv("user_basic_info.csv",names=["uId", "gender", "city", "prodName","ramCa
2 user_basic_info.tail(10)
```

Out[6]:

|         | uld     | gender | city   | prodName | ramCapacity | ramLeftRation | romCapacity | romLeftRation | color | fontSize | ct      | c        |
|---------|---------|--------|--------|----------|-------------|---------------|-------------|---------------|-------|----------|---------|----------|
| 2512490 | 3512472 | 0      | c00145 | p00147   | 4.0         | 0.32          | 64.0        | 0.55          | 幻夜黑   | 1.00     | wifi    | China_Te |
| 2512491 | 3512473 | 0      | c0066  | p009     | 3.0         | 0.36          | 64.0        | 0.66          | 荣耀金   | 1.15     | 4g      | China_Te |
| 2512492 | 3512474 | 1      | c00145 | p00169   | 3.0         | 0.52          | 32.0        | 0.12          | 皓月银   | 1.15     | 3g#wifi | China_U  |
| 2512493 | 3512477 | 0      | c00106 | p0044    | 2.0         | 0.27          | 16.0        | 0.24          | 太空银   | 1.00     | 4g#wifi | China_M  |
| 2512494 | 3512478 | 0      | c00281 | p00155   | 6.0         | 0.32          | 64.0        | 0.46          | 备件颜色  | 1.00     | 4g#wifi | China_Te |
| 2512495 | 3512489 | 0      | c00192 | p00102   | 6.0         | 0.36          | 128.0       | 0.74          | 翡冷翠   | 1.00     | 4g#wifi | China_U  |
| 2512496 | 3512491 | 0      | c007   | p00223   | 3.0         | 0.35          | 32.0        | 0.16          | 晨曦金   | NaN      | NaN     | China_Te |
| 2512497 | 3512492 | 0      | c0038  | p0016    | 6.0         | 0.33          | 64.0        | 0.20          | 琥珀金   | 1.00     | 4g#wifi | China_U  |
| 2512498 | 3512494 | 0      | c0043  | p0016    | 4.0         | 0.34          | 64.0        | 0.41          | 海鸥灰   | NaN      | 4g#wifi | China_M  |
| 2512499 | 3512495 | 0      | c00145 | p00141   | 6.0         | 0.26          | 68.0        | 0.54          | 银钻灰   | NaN      | wifi    | China_Te |

```
In [7]: 1 user_behavior_info=pd.read_csv("user_behavior_info.csv",names=["uId", "bootTimes", "AFuncTimes", "B
2 user_behavior_info.tail(10)
3
```

Out[7]:

|                | uld     | bootTimes | AFuncTimes | BFuncTimes | CFuncTimes | DFuncTimes | EFuncTimes | FFuncSum | GFuncSum |
|----------------|---------|-----------|------------|------------|------------|------------|------------|----------|----------|
| <b>2512490</b> | 3512491 | 5         | 0.0        | 0.00       | 0.00       | 0.00       | 0.00       | 0.0      | 0        |
| <b>2512491</b> | 3512492 | 24        | 0.0        | 0.00       | 0.07       | 0.43       | 0.07       | 2.0      | 154683   |
| <b>2512492</b> | 3512493 | 5         | 0.0        | 0.00       | 0.47       | 0.23       | 0.00       | 0.0      | 3193     |
| <b>2512493</b> | 3512494 | 4         | 0.0        | 0.07       | 1.47       | 0.00       | 0.00       | 0.0      | 1306     |
| <b>2512494</b> | 3512495 | 0         | 0.0        | 0.00       | 0.00       | 0.00       | 0.00       | 0.0      | 0        |
| <b>2512495</b> | 3512496 | 37        | 0.0        | 0.00       | 0.17       | 0.03       | 0.03       | 0.0      | 19431    |
| <b>2512496</b> | 3512497 | 40        | 0.0        | 0.00       | 0.00       | 0.67       | 0.07       | 0.0      | 947      |
| <b>2512497</b> | 3512498 | 2         | 0.0        | 0.00       | 0.00       | 0.00       | 0.00       | 0.0      | 0        |
| <b>2512498</b> | 3512499 | 2         | 0.0        | 0.17       | 1.17       | 6.83       | 7.83       | 0.0      | 512      |
| <b>2512499</b> | 3512500 | 11        | 1.0        | 0.00       | 1.37       | 3.90       | 2.87       | 0.0      | 0        |

```
In [8]: 1 user_app_activated=pd.read_csv("user_app_activated.csv",names=["uId", "appId"])
        2 user_app_activated.tail(10)
```

Out[8]:

|                | uld     | appld   |
|----------------|---------|---|
| <b>2512490</b> | 3512461 | a00102782#a00104275#a00109386#a00109824#a00118... |
| <b>2512491</b> | 3512464 | a00102720#a00109386#a00145168#a0015685#a001704... |
| <b>2512492</b> | 3512466 | a00102939#a00109386#a00134746#a00135785#a00150... |
| <b>2512493</b> | 3512475 | a00109386#a00144187#a00157201#a0021880#a002447... |
| <b>2512494</b> | 3512476 | a00102448#a00109386#a00144187#a00157058#a00157... |
| <b>2512495</b> | 3512477 | a00109386#a00121683#a00158614#a00187480#a00225... |
| <b>2512496</b> | 3512479 | a00109386#a0011894#a00145168#a00149248#a001586... |
| <b>2512497</b> | 3512483 | a00109386#a00135785#a0017008#a00170432#a001874... |
| <b>2512498</b> | 3512484 | a00102895#a00109386#a00158614#a00170432#a00212... |
| <b>2512499</b> | 3512491 | a00109386#a0012409#a00142198#a0014473#a0020705... |

```
In [9]: 1 app_info=pd.read_csv("app_info.csv",names=["appId", "category"])
        2 app_info.tail(10)
```

Out[9]:

|        | appld     | category |
|--------|-----------|----------|
| 188854 | a00488097 | 实用工具     |
| 188855 | a00488109 | 便捷生活     |
| 188856 | a00488109 | 实用工具     |
| 188857 | a00488110 | 便捷生活     |
| 188858 | a00488111 | 休闲益智     |
| 188859 | a00488113 | 经营策略     |
| 188860 | a00488114 | 便捷生活     |
| 188861 | a00488115 | 实用工具     |
| 188862 | a00488119 | 便捷生活     |
| 188863 | a00488122 | 实用工具     |

```
In [10]: 1 merged_inner = pd.merge(left=age_train,right=user_basic_info, left_on='uId', right_on='uId')
        2 merged_inner.tail()
```

Out[10]:

|         | uld     | age_group | gender | city   | prodName | ramCapacity | ramLeftRation | romCapacity | romLeftRation | color | fontSize |
|---------|---------|-----------|--------|--------|----------|-------------|---------------|-------------|---------------|-------|----------|
| 2009995 | 3512496 | 1         | 0      | c00245 | p00210   | 4.0         | 0.55          | 64.0        | 0.33          | 银色    | 1.0 4g#v |
| 2009996 | 3512497 | 6         | 0      | c00136 | p00208   | 3.0         | 0.41          | 32.0        | 0.76          | 金色    | 1.3 Nε   |
| 2009997 | 3512498 | 2         | 0      | c00176 | p0045    | 2.0         | 0.29          | 16.0        | 0.28          | 银色    | NaN Nε   |
| 2009998 | 3512499 | 2         | 0      | c00126 | p0022    | 8.0         | 0.17          | 128.0       | 0.84          | 渐变蓝   | 1.0 4g#v |
| 2009999 | 3512500 | 3         | 0      | c00229 | p00110   | 4.0         | 0.29          | 64.0        | 0.06          | 琥珀金   | 1.0 v    |

```
In [11]: 1 user_behavior_info.head()
```

```
Out[11]:
```

|   | uld     | bootTimes | AFuncTimes | BFuncTimes | CFuncTimes | DFuncTimes | EFuncTimes | FFuncSum | GFuncSum |
|---|---------|-----------|------------|------------|------------|------------|------------|----------|----------|
| 0 | 1000001 | 108       | 0.0        | 0.00       | 1.00       | 0.07       | 0.0        | 0.0      | 3319     |
| 1 | 1000002 | 14        | 0.0        | 0.17       | 4.93       | 1.23       | 3.9        | 1.0      | 245      |
| 2 | 1000003 | 13        | 1.0        | 0.00       | 7.73       | 3.00       | 1.7        | 0.0      | 5987     |
| 3 | 1000004 | 57        | 0.0        | 0.03       | 1.37       | 0.63       | 0.0        | 0.0      | 7460     |
| 4 | 1000005 | 0         | 0.0        | 0.00       | 0.00       | 0.00       | 0.0        | 0.0      | 0        |

```
In [12]: 1 to_predict=pd.merge(left=merged_inner,right=user_behavior_info, left_on='uId', right_on='uId')
          2 to_predict.tail()
```

```
Out[12]:
```

|         | uld     | age_group | gender | city   | prodName | ramCapacity | ramLeftRation | romCapacity | romLeftRation | color | ... | car         |
|---------|---------|-----------|--------|--------|----------|-------------|---------------|-------------|---------------|-------|-----|-------------|
| 2009995 | 3512496 | 1         | 0      | c00245 | p00210   | 4.0         | 0.55          | 64.0        | 0.33          | 银色    | ... | China_Telec |
| 2009996 | 3512497 | 6         | 0      | c00136 | p00208   | 3.0         | 0.41          | 32.0        | 0.76          | 金色    | ... | China_Unic  |
| 2009997 | 3512498 | 2         | 0      | c00176 | p0045    | 2.0         | 0.29          | 16.0        | 0.28          | 银色    | ... | China_Mo    |
| 2009998 | 3512499 | 2         | 0      | c00126 | p0022    | 8.0         | 0.17          | 128.0       | 0.84          | 渐变蓝   | ... | China_Unic  |
| 2009999 | 3512500 | 3         | 0      | c00229 | p00110   | 4.0         | 0.29          | 64.0        | 0.06          | 琥珀金   | ... | China_Mo    |

5 rows × 22 columns

```
In [13]: 1 to_predict['carrier'].unique()
```

```
Out[13]: array(['China_Mobile', 'China_Telecom', 'China_Unicom', 'othercp'],
              dtype=object)
```

```
In [14]: 1 to_predict.shape
```

```
Out[14]: (2010000, 22)
```

```
In [15]: 1 to_predict=to_predict.dropna()  
        2 to_predict.reset_index()  
        3 to_predict.shape
```

```
Out[15]: (1370640, 22)
```

## Feature engineering

```
In [16]: 1 X_labels=["gender", "city", "prodName", "ramCapacity", "ramLeftRation", "romCapacity", "romLeftRation", "color", "carrier", "os", "bootTime"]
2 y_label='age_group'
3
4 from sklearn.preprocessing import LabelEncoder
5
6 lb_make = LabelEncoder()
7 to_predict["carrier"] = lb_make.fit_transform(to_predict["carrier"])
8 to_predict["color"] = lb_make.fit_transform(to_predict["color"])
9 to_predict["ct"] = lb_make.fit_transform(to_predict["ct"])
10 to_predict["prodName"] = lb_make.fit_transform(to_predict["prodName"])
11 to_predict["city"] = lb_make.fit_transform(to_predict["city"])
12 to_predict.head(11)
13
```

Out[16]:

|    | uld     | age_group | gender | city | prodName | ramCapacity | ramLeftRation | romCapacity | romLeftRation | color | ... | carrier | os  | bootTi |
|----|---------|-----------|--------|------|----------|-------------|---------------|-------------|---------------|-------|-----|---------|-----|--------|
| 0  | 1000001 | 4         | 0      | 48   | 50       | 3.0         | 0.43          | 32.0        | 0.46          | 70    | ... | 0       | 8.0 |        |
| 2  | 1000015 | 5         | 1      | 223  | 50       | 3.0         | 0.34          | 32.0        | 0.06          | 70    | ... | 1       | 8.0 |        |
| 4  | 1000023 | 2         | 1      | 288  | 105      | 2.0         | 0.34          | 16.0        | 0.06          | 103   | ... | 1       | 7.0 |        |
| 5  | 1000025 | 4         | 0      | 311  | 52       | 4.0         | 0.31          | 64.0        | 0.20          | 65    | ... | 1       | 8.0 |        |
| 6  | 1000029 | 4         | 0      | 346  | 41       | 6.0         | 0.20          | 68.0        | 0.27          | 31    | ... | 2       | 9.0 |        |
| 7  | 1000035 | 2         | 0      | 234  | 41       | 6.0         | 0.26          | 68.0        | 0.64          | 12    | ... | 1       | 9.0 |        |
| 9  | 1000038 | 3         | 0      | 329  | 10       | 4.0         | 0.29          | 64.0        | 0.16          | 42    | ... | 0       | 7.0 |        |
| 10 | 1000040 | 3         | 0      | 247  | 32       | 6.0         | 0.20          | 128.0       | 0.52          | 86    | ... | 1       | 9.0 |        |
| 11 | 1000044 | 4         | 0      | 323  | 43       | 6.0         | 0.22          | 137.0       | 0.30          | 2     | ... | 2       | 9.0 |        |
| 12 | 1000046 | 3         | 0      | 131  | 106      | 4.0         | 0.49          | 64.0        | 0.01          | 7     | ... | 0       | 7.0 |        |
| 13 | 1000051 | 3         | 1      | 28   | 45       | 6.0         | 0.27          | 128.0       | 0.17          | 115   | ... | 0       | 8.0 |        |

11 rows × 22 columns

```
In [17]: 1 to_predict.shape
```

Out[17]: (1370640, 22)



```
In [18]: 1 to_predict['age_group'].value_counts()
```

```
Out[18]: 3    406461
         4    345040
         2    258685
         5    214086
         6    109579
         1     36789
         Name: age_group, dtype: int64
```

```
In [66]: 1 y_data=to_predict[y_label]
         2 X_data=to_predict[X_labels]
         3 # del to_predict['age_group']
         4 # to_predict.tail()
         5 print(X_data.shape)
         6 y_data.shape
```

```
(1370640, 20)
```

```
Out[66]: (1370640,)
```

```
In [73]: 1 from sklearn import preprocessing
         2 X = X_data.values #returns a numpy array
         3 min_max_scaler = preprocessing.MinMaxScaler()
         4 X_scaled = min_max_scaler.fit_transform(X)
         5 X_scaled
         6 # df = pandas.DataFrame(x_scaled)
```

```
Out[73]: array([[0.00000000e+00, 1.36363636e-01, 3.96825397e-01, ...,
                 5.00000000e-01, 0.00000000e+00, 7.98260115e-05],
                [1.00000000e+00, 6.33522727e-01, 3.96825397e-01, ...,
                 5.00000000e-01, 0.00000000e+00, 5.26264826e-04],
                [1.00000000e+00, 8.18181818e-01, 8.33333333e-01, ...,
                 5.00000000e-01, 0.00000000e+00, 0.00000000e+00],
                ...,
                [0.00000000e+00, 4.43181818e-01, 5.39682540e-01, ...,
                 5.00015000e-01, 0.00000000e+00, 4.67339328e-04],
                [0.00000000e+00, 7.95454545e-02, 5.95238095e-01, ...,
                 5.03915000e-01, 0.00000000e+00, 1.23142265e-05],
                [0.00000000e+00, 3.94886364e-01, 8.73015873e-02, ...,
                 5.01435000e-01, 0.00000000e+00, 0.00000000e+00]])
```

```
In [67]: 1 y=y_data.values
          2 y
```

```
Out[67]: array([4, 5, 2, ..., 1, 2, 3])
```

## Model Training

```
In [75]: 1 from sklearn.model_selection import train_test_split
          2 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=11)
          3
          4 # for experimental purpose, we only use 10% of data to train within a perceivable period of time
          5 # X_train, X_test, y_train, y_test = train_test_split(X_test, y_test, test_size=0.2, random_state=1
          6 # X_train, X_test, y_train, y_test = train_test_split(X_test, y_test, test_size=0.2, random_state=1
          7
          8 print(X_train.shape)
          9 print(y_train.shape)
         10 print("-----")
         11 print(X_test.shape)
         12 print(y_test.shape)
```

```
(1096512, 20)
```

```
(1096512,)
```

```
-----
```

```
(274128, 20)
```

```
(274128,)
```

```
In [77]: 1 from sklearn.naive_bayes import GaussianNB
2 ml_m = GaussianNB()
3 %time ml_m.fit(X_train, y_train)
4 print('Accuracy of GaussianNB classifier on Test Set: {:.4f}'
5       .format(ml_m.score(X_test, y_test)))
6
7 predicted_test = ml_m.predict(X_test)
8 from sklearn.metrics import classification_report
9 print(classification_report(y_test, predicted_test))
```

CPU times: user 566 ms, sys: 262 ms, total: 829 ms

Wall time: 891 ms

Accuracy of GaussianNB classifier on Test Set: 0.2412

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1            | 0.04      | 0.41   | 0.08     | 7164    |
| 2            | 0.28      | 0.53   | 0.37     | 51697   |
| 3            | 0.35      | 0.16   | 0.22     | 81683   |
| 4            | 0.38      | 0.12   | 0.19     | 69078   |
| 5            | 0.29      | 0.03   | 0.05     | 42510   |
| 6            | 0.28      | 0.59   | 0.38     | 21996   |
| micro avg    | 0.24      | 0.24   | 0.24     | 274128  |
| macro avg    | 0.27      | 0.31   | 0.21     | 274128  |
| weighted avg | 0.32      | 0.24   | 0.22     | 274128  |

```
In [82]: 1 from sklearn.svm import LinearSVC
2 lsvc_pri_model = LinearSVC(C=0.1)
3
4 %time lsvc_pri_model.fit(X_train, y_train)
5 print('Accuracy of LinearSVC classifier on Test Set: {:.4f}'
6       .format(lsvc_pri_model.score(X_test, y_test)))
7
8 predicted_test = lsvc_pri_model.predict(X_test)
9 from sklearn.metrics import classification_report
10 print(classification_report(y_test, predicted_test))
```

CPU times: user 1min 27s, sys: 1.52 s, total: 1min 29s

Wall time: 1min 37s

Accuracy of LinearSVC classifier on Test Set: 0.3390

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1            | 0.00      | 0.00   | 0.00     | 7164    |
| 2            | 0.36      | 0.05   | 0.08     | 51697   |
| 3            | 0.34      | 0.80   | 0.48     | 81683   |
| 4            | 0.34      | 0.16   | 0.22     | 69078   |
| 5            | 0.31      | 0.30   | 0.30     | 42510   |
| 6            | 0.42      | 0.06   | 0.10     | 21996   |
| micro avg    | 0.34      | 0.34   | 0.34     | 274128  |
| macro avg    | 0.30      | 0.23   | 0.20     | 274128  |
| weighted avg | 0.34      | 0.34   | 0.27     | 274128  |

/Library/Frameworks/Python.framework/Versions/3.7/lib/python3.7/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.

'precision', 'predicted', average, warn\_for)

## Prediction for test set and submission

```

In [89]: 1 merged_inner_test = pd.merge(left=age_test,right=user_basic_info, left_on='uId', right_on='uId')
2 to_predict_onetest=pd.merge(left=merged_inner_test,right=user_behavior_info, left_on='uId', right_on='uId')
3 # to_predict_onetest.tail()
4 print(to_predict_onetest.shape)
5 to_predict_onetest=to_predict_onetest.dropna()
6 print(to_predict_onetest.shape)
7
8 from sklearn.preprocessing import LabelEncoder
9
10 lb_make = LabelEncoder()
11 to_predict_onetest["carrier"] = lb_make.fit_transform(to_predict_onetest["carrier"])
12 to_predict_onetest["color"] = lb_make.fit_transform(to_predict_onetest["color"])
13 to_predict_onetest["ct"] = lb_make.fit_transform(to_predict_onetest["ct"])
14 to_predict_onetest["prodName"] = lb_make.fit_transform(to_predict_onetest["prodName"])
15 to_predict_onetest["city"] = lb_make.fit_transform(to_predict_onetest["city"])
16 to_predict_onetest.head(11)

```

(502500, 21)

(342926, 21)

Out[89]:

|    | uld     | gender | city | prodName | ramCapacity | ramLeftRation | romCapacity | romLeftRation | color | fontSize | ... | carrier | os  | bootTime |
|----|---------|--------|------|----------|-------------|---------------|-------------|---------------|-------|----------|-----|---------|-----|----------|
| 1  | 1000003 | 0      | 219  | 4        | 8.0         | 0.35          | 128.0       | 0.60          | 78    | 1.00000  | ... | 0       | 9.0 | 1        |
| 2  | 1000004 | 1      | 195  | 22       | 3.0         | 0.36          | 32.0        | 0.09          | 18    | 1.00000  | ... | 0       | 8.0 | 5        |
| 5  | 1000007 | 1      | 310  | 120      | 6.0         | 0.20          | 137.0       | 0.50          | 12    | 1.00000  | ... | 1       | 9.0 | 1        |
| 6  | 1000008 | 1      | 294  | 39       | 6.0         | 0.36          | 128.0       | 0.47          | 84    | 1.15000  | ... | 1       | 8.0 |          |
| 7  | 1000009 | 0      | 332  | 39       | 6.0         | 0.39          | 64.0        | 0.23          | 97    | 1.15000  | ... | 0       | 8.0 | 2        |
| 9  | 1000012 | 0      | 198  | 48       | 4.0         | 0.31          | 64.0        | 0.39          | 65    | 1.30001  | ... | 0       | 8.0 | 9        |
| 10 | 1000013 | 0      | 131  | 38       | 4.0         | 0.49          | 64.0        | 0.48          | 97    | 1.00000  | ... | 0       | 8.0 | 2        |
| 11 | 1000014 | 1      | 310  | 2        | 6.0         | 0.49          | 128.0       | 0.58          | 31    | 1.00000  | ... | 0       | 9.0 | 1        |
| 12 | 1000016 | 0      | 343  | 89       | 3.0         | 0.41          | 32.0        | 0.46          | 99    | 1.00000  | ... | 2       | 7.0 | 1        |
| 13 | 1000017 | 1      | 293  | 11       | 4.0         | 0.29          | 64.0        | 0.04          | 65    | 1.15000  | ... | 0       | 8.0 | 2        |
| 15 | 1000020 | 0      | 223  | 80       | 4.0         | 0.38          | 64.0        | 0.09          | 28    | 1.00000  | ... | 2       | 8.0 | 1        |

11 rows × 21 columns

```

In [90]: 1 test_labels=["gender", "city","prodName","ramCapacity","ramLeftRation","romCapacity","romLeftRation"
2
3 test_data=to_predict_ontest[test_labels]
4
5 from sklearn import preprocessing
6 test = test_data.values #returns a numpy array
7
8 min_max_scaler = preprocessing.MinMaxScaler()
9 test_scaled = min_max_scaler.fit_transform(test)
10 test_scaled
11
12 # test_scaled=test

```

```

Out[90]: array([[0.00000000e+00, 6.23931624e-01, 3.27868852e-02, ...,
3.79937888e-01, 0.00000000e+00, 3.65952114e-04],
[1.00000000e+00, 5.55555556e-01, 1.80327869e-01, ...,
3.78881988e-01, 0.00000000e+00, 4.55988436e-04],
[1.00000000e+00, 8.83190883e-01, 9.83606557e-01, ...,
3.79670807e-01, 0.00000000e+00, 0.00000000e+00],
...,
[0.00000000e+00, 9.37321937e-01, 6.06557377e-01, ...,
3.79254658e-01, 5.00000000e-02, 5.80437827e-04],
[1.00000000e+00, 4.13105413e-01, 1.88524590e-01, ...,
3.78881988e-01, 0.00000000e+00, 0.00000000e+00],
[0.00000000e+00, 3.10541311e-01, 1.80327869e-01, ...,
3.80726708e-01, 0.00000000e+00, 8.04397831e-05]])

```

```

In [91]: 1 predicted = lsvc_pri_model.predict(test_scaled)
2 predicted

```

```

Out[91]: array([5, 2, 5, ..., 5, 5, 2])

```

```

In [92]: 1 print(to_predict_ontest['uId'].shape)
2 print(len(predicted))

```

```

(342926,)
342926

```

```

In [93]: 1 sub_df=pd.DataFrame(to_predict_ontest['uId']).reset_index(drop=True)
2 sub_df['label']=predicted

```

```
In [94]: 1 sub_df.columns=['id','label']  
2 sub_df.to_csv("submission.csv",index=False)  
3 print("submission.csv is saved sucessfully!")
```

submission.csv is saved sucessfully!

```
In [ ]: 1
```